



# Machine learning workshops for material scientists

## Lecture 6: Supervised linear model with Python

October 31, 2022



**Sira Sriswasdi, PhD**

- Research Affairs
- Center of Excellence in Computational Molecular Biology (CMB)
- Center for Artificial Intelligence in Medicine (CU-AIM)

# Today's goals



- Quick recap of linear models
- Get to know Python library
- Practice on Colab




# Linear and logistic regression

# Linear and logistic regression key points



- Predictions are made via linear combination of input features
  - $\hat{y} = b_0 + b_1x_1 + \dots + b_nx_n$
  - $\log\left(\frac{\hat{y}}{1-\hat{y}}\right) = b_0 + b_1x_1 + \dots + b_nx_n$
- Primary objectives:
  - MSE:  $\frac{1}{n}\sum(y - \hat{y})^2$
  - Log-Likelihood, Cross-Entropy:  $\sum y \log(\hat{y}) + (1 - y) \log(1 - \hat{y})$
- Regularization
  - LASSO: objective +  $\alpha \sum |b_i|$  or Ridge: Objective +  $\alpha \sum b_i^2$

# Linear regression



<https://sixsigmadsi.com/glossary/simple-linear-regression/>

- $\hat{y} = b_0 + b_1x_1 + \dots + b_nx_n$
- Minimize MSE:  $\frac{1}{n}\sum(y - \hat{y})^2$
- MSE objective is related to assumption that errors are normally distributed and centered at zero
- Data transformation (e.g., log) can be beneficial

# Linear regression in Python

## sklearn.linear\_model.LinearRegression

```
class sklearn.linear_model.LinearRegression(*, fit_intercept=True, normalize='deprecated', copy_X=True, n_jobs=None, positive=False) \[source\]
```

```
from sklearn.linear_model import LinearRegression
```

```
linear = LinearRegression().fit(input_data, label)
```

```
prediction = linear.predict(new_data)
```

```
## array of fitted coefficients ( $b_1, b_2, \dots, b_n$ ) for input features
```

```
linear.coef_
```

```
## y-intercept ( $b_0$ )
```

```
linear.intercept_
```

# Sample weighting

```
fit(X, y, sample_weight=None)
```

Fit linear model.

**Parameters::**

**X** : {array-like, sparse matrix} of shape (n\_samples, n\_features)

Training data.

**y** : array-like of shape (n\_samples,) or (n\_samples, n\_targets)

Target values. Will be cast to X's dtype if necessary.

**sample\_weight** : array-like of shape (n\_samples,), default=None

Individual weights for each sample.

- Add weight to the calculation of objective:  $\sum w_i (y_i - \hat{y}_i)^2$
- Large weight  $\rightarrow$  focus on fitting those data points

# Ridge and LASSO in Python



## `sklearn.linear_model.Ridge`

```
class sklearn.linear_model.Ridge(alpha=1.0, *, fit_intercept=True, normalize='deprecated', copy_X=True, max_iter=None, tol=0.001, solver='auto', positive=False, random_state=None) \[source\]
```

## `sklearn.linear_model.Lasso`

```
class sklearn.linear_model.Lasso(alpha=1.0, *, fit_intercept=True, normalize='deprecated', precompute=False, copy_X=True, max_iter=1000, tol=0.0001, warm_start=False, positive=False, random_state=None, selection='cyclic') \[source\]
```

- **alpha** = regularization strength
  - Larger alpha → try to keep coefficients small
  - Small alpha → focus more on prediction performance



# Ridge and LASSO usage



```
from sklearn.linear_model import Ridge, Lasso
```

```
## Ridge model
```

```
ridge = Ridge(alpha = 10, random_state = 25)  
ridge.fit(input_data, label)
```

```
## LASSO model
```

```
lasso = Lasso(alpha = 10, random_state = 25)  
lasso.fit(input_data, label)
```

```
## make predictions
```

```
prediction = ridge.predict(new_data)
```

# Logistic regression




Image from Wikipedia

- Transform linear regression output into probability via log-odd function
- $\log\left(\frac{\hat{y}}{1-\hat{y}}\right) = b_0 + b_1x_1 + \dots + b_nx_n$
- Minimize Cross-Entropy:
  - $\sum y \log(\hat{y}) + (1 - y) \log(1 - \hat{y})$
- Data transformation (e.g., log) can be beneficial
- Well-calibrated (interpretable output)

# Logistic regression in Python



## `sklearn.linear_model.LogisticRegression`

```
class sklearn.linear_model.LogisticRegression(penalty='l2', *, dual=False, tol=0.0001, C=1.0, fit_intercept=True,
intercept_scaling=1, class_weight=None, random_state=None, solver='lbfgs', max_iter=100, multi_class='auto', verbose=0,
warm_start=False, n_jobs=None, l1_ratio=None)
```

- **penalty** = type of regularization (l1 = LASSO, l2 = Ridge)
- **C** = inverse of regularization strength (1 / alpha)
- **solver** = optimization algorithm to be used
  - Not all support LASSO regularization → change to “liblinear”

# Logistic regression usage



```
from sklearn.linear_model import LogisticRegression

logreg = LogisticRegression(penalty = 'l2', C = 0.01, random_state = 25)
logreg.fit(input_data, label)

## coefficients
logreg.coef_
logreg.intercept_


## predict classes
pred_class = logreg.predict(new_data)

## predict raw probabilities
pred_prob = logreg.predict_proba(new_data)
```



# Support vector machine

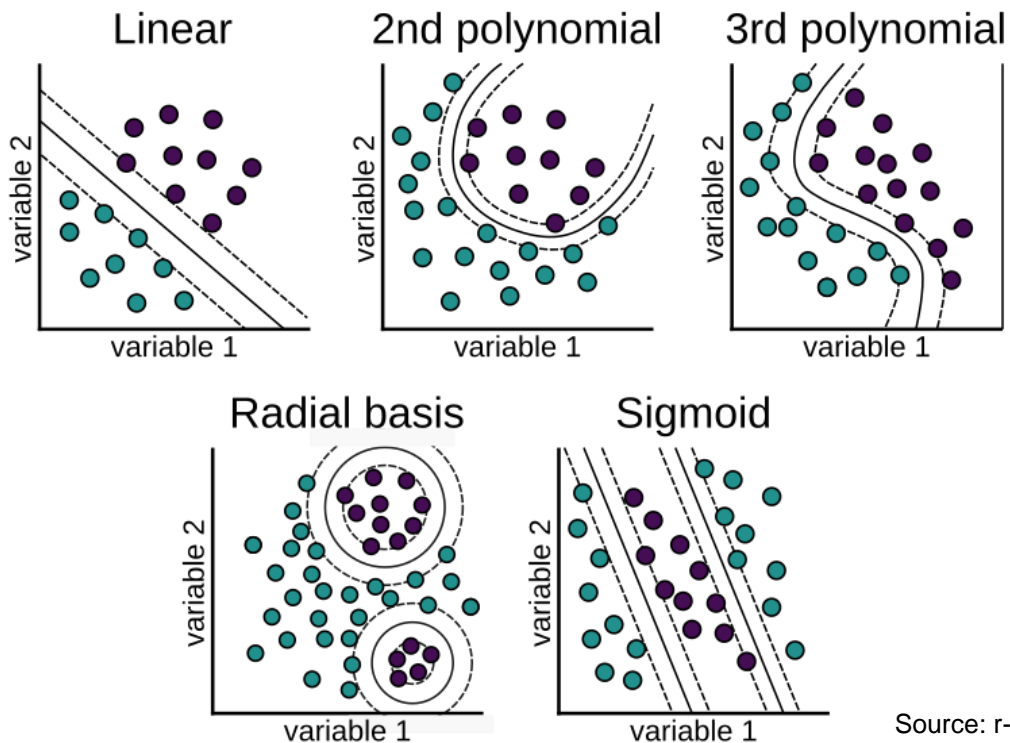
# Support vector machine key points



Source: towarddatasciences.com


- Hinge loss
  - Penalize correctly classified data points that lie too close to the decision boundary
  - Do not penalize small regression error
- Generalize to non-linear feature engineering via kernel
  - Radial basis
  - Polynomial
- Does not output class probabilities

# Kernels




- Radial basis operates like nearest neighbor algorithm
  - Require large training data to be effective
- Polynomial can work well when there are interactions between features
  - Terms  $x_i x_j$  are produced

# Support vector classifier versus regressor



Source: wikipedia.com



Rosenbaum, L. et al. J of Cheminformatics 5:33 (2013)



# Support vector classifier in Python

## `sklearn.svm.LinearSVC`

```
class sklearn.svm.LinearSVC(penalty='l2', loss='squared_hinge', *, dual=True, tol=0.0001, C=1.0, multi_class='ovr',  
fit_intercept=True, intercept_scaling=1, class_weight=None, verbose=0, random_state=None, max_iter=1000)
```

## `sklearn.svm.SVC`

```
class sklearn.svm.SVC(*, C=1.0, kernel='rbf', degree=3, gamma='scale', coef0=0.0, shrinking=True, probability=False, tol=0.001,  
cache_size=200, class_weight=None, verbose=False, max_iter=-1, decision_function_shape='ovr', break_ties=False,  
random_state=None)
```

[\[source\]](#)

- **C** = inverse of regularization strength
- **kernel**, **degree** = set kernel for non-linear SVM
- **probability** = make the model estimate probability (empirically from cross-validation), enable **predict\_proba**

# Support vector classifier usage



```
from sklearn.svm import LinearSVC, SVC
```

```
## linear kernel
```

```
linearsvc = LinearSVC(C = 0.01, random_state = 25)
```

```
linearsvc.fit(input_data, label)
```

```
prediction = linearsvc.predict(new_data)
```

```
linearsvc.coef_
```

```
## non-linear kernel
```

```
svc = SVC(C = 0.01, kernel = 'poly', degree = 2, random_state = 25)
```

```
svc.fit(input_data, label)
```

```
prediction = svc.predict(new_data)
```

```
## get all support vectors (data points near decision boundary)
```

```
svc.support_vectors_
```

# Support vector regressor in Python

## `sklearn.svm.LinearSVR`

```
class sklearn.svm.LinearSVR(*, epsilon=0.0, tol=0.0001, C=1.0, loss='epsilon_insensitive', fit_intercept=True, intercept_scaling=1.0, dual=True, verbose=0, random_state=None, max_iter=1000)
```

## `sklearn.svm.SVR`

```
class sklearn.svm.SVR(*, kernel='rbf', degree=3, gamma='scale', coef0=0.0, tol=0.001, C=1.0, epsilon=0.1, shrinking=True, cache_size=200, verbose=False, max_iter=-1)
```

- **C** = inverse of regularization strength
- **kernel**, **degree** = set kernel for non-linear SVM
- **epsilon** = maximum regression error to not penalize

# Support vector regressor usage



```
from sklearn.svm import LinearSVR, SVR
```

```
## linear kernel
```

```
linearsvr = LinearSVR(C = 0.01, epsilon = 0.5, random_state = 25)
```

```
linearsvr.fit(input_data, label)
```

```
prediction = linearsvr.predict(new_data)
```

```
linearsvr.coef_
```

```
## non-linear kernel
```

```
svr = SVR(C = 0.01, kernel = 'poly', degree = 2, epsilon = 0.5, random_state = 25)
```

```
svr.fit(input_data, label)
```

```
prediction = svr.predict(new_data)
```

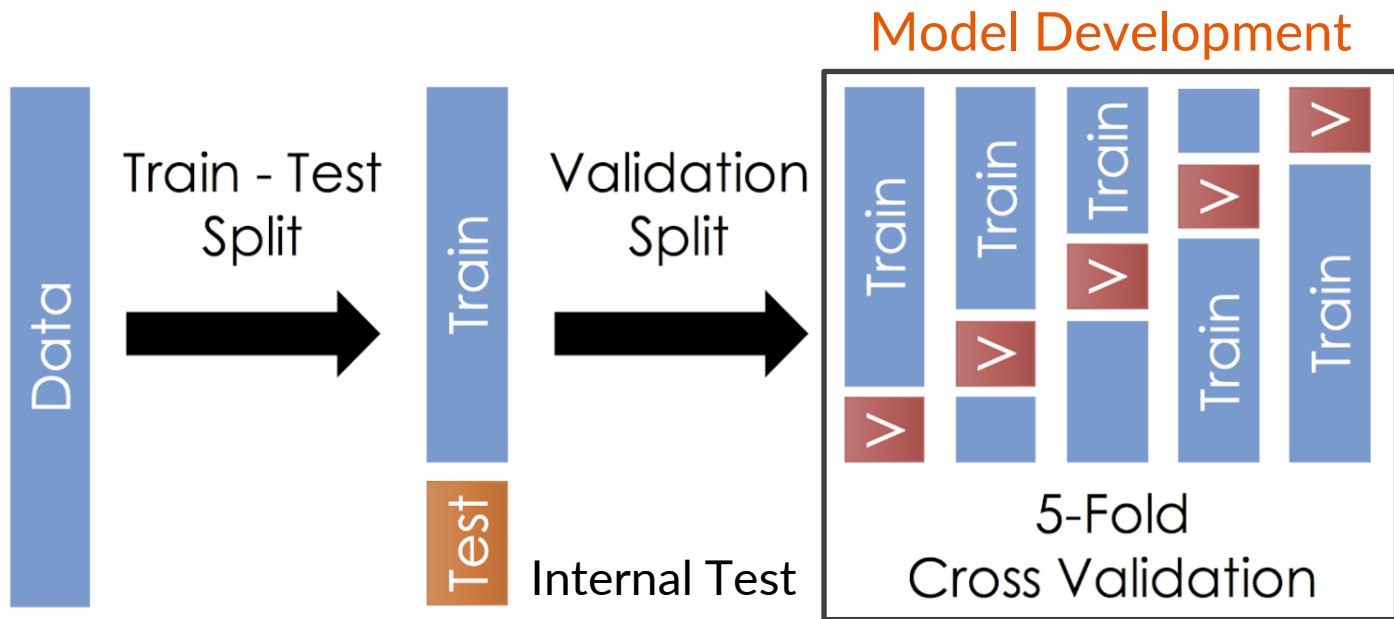
```
## get all support vectors (data points near decision boundary)
```

```
svr.support_vectors_
```



# Data splitting

# Train-Val-Test



# Data splitting tips



- **Test or no test:** can you dedicate enough sample to test set?
- Label distribution on validation and test set
  - Stratified versus 1:1
- **Cross-validation** versus **bootstrapping**
  - CV ensures that all samples will be used equally
  - Bootstrapping provides more control on label distribution
- Monitor performances on each split
  - Mean and SD

# Data splitting examples



- **Example 1:** 223 negative, 77 positive
  - **Test:** 31 negative, 27 positive
  - **Validation:** 25 negative, 25 positive
  - **Training:** 167 negative, 25 positive
  - Repeat training-validation split using bootstrapping
- **Example 2:** 48 negative, 23 positive
  - 2-fold cross-validation: 24 negative, 11 positive
  - Perform only logistic regression model
  - Limited to discussion of feature importance



# Data splitting in Python



## `sklearn.model_selection.train_test_split`

```
sklearn.model_selection.train_test_split(*arrays, test_size=None, train_size=None, random_state=None, shuffle=True, stratify=None)
```

[\[source\]](#)

```
from sklearn.model_selection import train_test_split
```

```
## random split
```

```
X_train, X_test, y_train, y_test = train_test_split(input_data, label, test_size = 0.25, random_state = 25)
```

```
## stratified split
```

```
X_train, X_test, y_train, y_test = train_test_split(input_data, label, test_size = 0.25, stratify = label, random_state = 25)
```

# Cross-validation in Python



`sklearn.model_selection.KFold`

```
class sklearn.model_selection.KFold(n_splits=5, *, shuffle=False, random_state=None)
```

`sklearn.model_selection.StratifiedKFold`

```
class sklearn.model_selection.StratifiedKFold(n_splits=5, *, shuffle=False, random_state=None)
```

- `n_splits` = number of folds
- Fixed random state and set `shuffle = True` (otherwise will get consecutive)
- Stratification will always be done based on the label (y)

# Cross-validation usage



```
from sklearn.model_selection import StratifiedKFold

skf = StratifiedKFold(n_splits = 10, shuffle = True, random_state = 25)

## a fitted PCA can be used to transform any new data
for train_index, test_index in skf.split(input_data, label):
    X_train = input_data[train_index, :]
    y_train = label[train_index]

    X_test = input_data[test_index, :]
    y_test = label[test_index]

## model development
```

# Bootstrapping in Python



## `sklearn.model_selection.ShuffleSplit`

```
class sklearn.model_selection.ShuffleSplit(n_splits=10, *, test_size=None, train_size=None, random_state=None)
```

## `sklearn.model_selection.StratifiedShuffleSplit`

```
class sklearn.model_selection.StratifiedShuffleSplit(n_splits=10, *, test_size=None, train_size=None, random_state=None)
```

- `n_splits` = number of bootstrap samples
- Controllable sizes of both train and test splits
- Stratification will be done based on the label (y)

# Bootstrapping usage



```
from sklearn.model_selection import StratifiedShuffleSplit

sss = StratifiedShuffleSplit(n_splits = 100, train_size = 0.5, test_size = 0.3,
random_state = 25)

## a fitted PCA can be used to transform any new data
for train_index, test_index in sss.split(input_data, label):
    X_train = input_data[train_index, :]
    y_train = label[train_index]

    X_test = input_data[test_index, :]
    y_test = label[test_index]

## model development
```



# Hyperparameter tuning

# Hyperparameter tuning tips



- Understand the impact of key parameters for each model family
  - Increase  $\alpha$  in regularization  $\alpha \sum b_i^2$  helps overfitting
  - Decrease  $\alpha$  will instead help underfitting
- No need to tune aggressively
  - $\alpha \in \{0.0001, 0.01, 1, 100\}$
  - $\text{max\_depth} \in \{5, 10, 20, \text{None}\}$
- Monitor multiple metrics: AUROC, average precision, etc.
- Look for pattern in the result
  - Large enough  $\alpha$  yields good performance

# Hyperparameter tuning in Python



## `sklearn.model_selection.GridSearchCV`

```
class sklearn.model_selection.GridSearchCV(estimator, param_grid, *, scoring=None, n_jobs=None, refit=True, cv=None, verbose=0, pre_dispatch='2*n_jobs', error_score=nan, return_train_score=False)
```

[\[source\]](#)

- **estimator** = number of directions/axes to produce
- **param\_grid** = dictionary of hyperparameters to tune
- **scoring** = metrics to keep track of
- **n\_jobs** = use multiple CPU to speed up the calculation
- **refit** = whether to train using the optimal hyperparameters
- **cv** = data splitting



# Hyperparameter tuning usage



```
## preparing data split, base model, hyperparameters
```

```
skf = StratifiedKFold(n_splits = 10, shuffle = True, random_state = 25)
logreg = LogisticRegression(solver = 'liblinear', max_iter = 1000, random_state = 25)
logreg_param = {'C': [0.0001, 0.01, 1, 100],
                'penalty': ['l1', 'l2']}
```

```
## define and fit GridSearchCV
```

```
grid_logreg = GridSearchCV(logreg, logreg_param, scoring = ['accuracy', 'roc_auc',
'average_precision'], n_jobs = 4, refit = 'average_precision', cv = skf)
```

```
grid_logreg.fit(input_data, label)
```

```
## dictionary containing full tuning results (performance scores, time spent, etc.)
```

```
grid_logreg.cv_results_
```

# GridSearchCV output



```
'param_kernel': masked_array(data = ['poly', 'poly', 'rbf', 'rbf'],
                              mask = [False False False False]...)
'param_gamma': masked_array(data = [-- -- 0.1 0.2],
                             mask = [ True  True False False]...),
'param_degree': masked_array(data = [2.0 3.0 -- --],
                              mask = [False False  True  True]...),
'split0_test_score' : [0.80, 0.70, 0.80, 0.93],
'split1_test_score' : [0.82, 0.50, 0.70, 0.78],
'mean_test_score'    : [0.81, 0.60, 0.75, 0.85],
'std_test_score'     : [0.01, 0.10, 0.05, 0.08],
'rank_test_score'    : [2, 4, 3, 1],
'split0_train_score' : [0.80, 0.92, 0.70, 0.93],
'split1_train_score' : [0.82, 0.55, 0.70, 0.87],
'mean_train_score'   : [0.81, 0.74, 0.70, 0.90],
'std_train_score'    : [0.01, 0.19, 0.00, 0.03],
'mean_fit_time'      : [0.73, 0.63, 0.43, 0.49],
'std_fit_time'       : [0.01, 0.02, 0.01, 0.01],
'mean_score_time'    : [0.01, 0.06, 0.04, 0.04],
'std_score_time'     : [0.00, 0.00, 0.00, 0.01],
'params'             : [{'kernel': 'poly', 'degree': 2}, ...],
```

# Any question?

